

MeriSkill Internship Program

Name: Sannidhya Das

Task 2: Diabetics Prediction

Description:

The objective of the dataset is to predict whether or not a patient has diabetes, based on certain diagnostic measurements included in the dataset. The datasets consists of several medical predictor variables and one target variable, Outcome. Predictor variables includes the number of pregnancies the patient has had, their BMI, insulin level, age, and so on.

Import libraries and Dataset

```
In [1]: import numpy as np
import pandas as pd
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn import svm
from sklearn.metrics import accuracy_score
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
```

importing the data

```
In [2]: data=pd.read_csv("diabetes.csv")
data
```

```
Out[2]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
0	6	148	72	35	0	33.6	0.627	50	1
1	1	85	66	29	0	26.6	0.351	31	0
2	8	183	64	0	0	23.3	0.672	32	1
3	1	89	66	23	94	28.1	0.167	21	0
4	0	137	40	35	168	43.1	2.288	33	1
...
763	10	101	76	48	180	32.9	0.171	63	0
764	2	122	70	27	0	36.8	0.340	27	0
765	5	121	72	23	112	26.2	0.245	30	0
766	1	126	60	0	0	30.1	0.349	47	1
767	1	93	70	31	0	30.4	0.315	23	0

768 rows × 9 columns

viewing top 5 rows

```
In [3]: data.head()
```

```
Out[3]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
0	6	148	72	35	0	33.6	0.627	50	1
1	1	85	66	29	0	26.6	0.351	31	0
2	8	183	64	0	0	23.3	0.672	32	1
3	1	89	66	23	94	28.1	0.167	21	0
4	0	137	40	35	168	43.1	2.288	33	1

viewing the dimensions of the dataset

```
In [4]: data.shape
```

Out[4]: (768, 9)

viewing dataframe summary

```
In [5]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Pregnancies            768 non-null   int64
1   Glucose                768 non-null   int64
2   BloodPressure          768 non-null   int64
3   SkinThickness          768 non-null   int64
4   Insulin                768 non-null   int64
5   BMI                   768 non-null   float64
6   DiabetesPedigreeFunction 768 non-null   float64
7   Age                   768 non-null   int64
8   Outcome                768 non-null   int64
dtypes: float64(2), int64(7)
memory usage: 54.1 KB
```

viewing descriptive statistics

```
In [6]: data.describe()
```

Out[6]:

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
count	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000
mean	3.845052	120.894531	69.105469	20.536458	79.799479	31.992578	0.471876	33.240885	0.348958
std	3.369578	31.972618	19.355807	15.952218	115.244002	7.884160	0.331329	11.760232	0.476951
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.078000	21.000000	0.000000
25%	1.000000	99.000000	62.000000	0.000000	0.000000	27.300000	0.243750	24.000000	0.000000
50%	3.000000	117.000000	72.000000	23.000000	30.500000	32.000000	0.372500	29.000000	0.000000
75%	6.000000	140.250000	80.000000	32.000000	127.250000	36.600000	0.626250	41.000000	1.000000
max	17.000000	199.000000	122.000000	99.000000	846.000000	67.100000	2.420000	81.000000	1.000000

```
In [7]: data.columns
```

Out[7]: Index(['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness', 'Insulin', 'BMI', 'DiabetesPedigreeFunction', 'Age', 'Outcome'], dtype='object')

Checking for null values

```
In [8]: data.isnull().sum()
```


Out[8]: Pregnancies 0
Glucose 0
BloodPressure 0
SkinThickness 0
Insulin 0
BMI 0
DiabetesPedigreeFunction 0
Age 0
Outcome 0
dtype: int64

Exploratory Data Analysis

In [9]: *#Pregnancies variable distn*

```
plt.figure(figsize=(15,7))
sns.histplot(data['Pregnancies'], facecolor='c',bins=100,stat='density')
plt.ylabel('Precentage')
plt.xlabel('Pregnancies')
```


Out[9]: Text(0.5, 0, 'Pregnancies')



In [10]: *#Glucose variable distn*

```
plt.figure(figsize=(15,7))
sns.histplot(data['Glucose'], facecolor='violet',bins=100,stat='density')
plt.ylabel('Precentage')
plt.xlabel('Glucose')
```

Out[10]: Text(0.5, 0, 'Glucose')




In [11]: *#Blood Pressure distn*

```
plt.figure(figsize=(15,7))
sns.histplot(data['BloodPressure'], facecolor='steelblue',bins=100,stat='density')
```

```
plt.ylabel('Precentage')
plt.xlabel('BloodPressure')
```


Out[11]: Text(0.5, 0, 'BloodPressure')



```
In [12]: #Skin Sickness distn

plt.figure(figsize=(15,7))
sns.histplot(data['SkinThickness'], facecolor='green',bins=100,stat='density')
plt.ylabel('Precentage')
plt.xlabel('SkinThickness')
```


Out[12]: Text(0.5, 0, 'SkinThickness')



```
In [13]: #BMI distn


plt.figure(figsize=(15,7))
sns.histplot(data['BMI'], facecolor='pink',bins=100,stat='density')
plt.ylabel('Precentage')
plt.xlabel('BMI')
```

Out[13]: Text(0.5, 0, 'BMI')




```
In [14]: # Diabetes Pedigree Function EDA
plt.figure(figsize=(15,7))
sns.histplot(data['DiabetesPedigreeFunction'], facecolor='seagreen',bins=100,stat='density')
plt.ylabel('Percentage')
plt.xlabel('DiabetesPedigreeFunction')
```

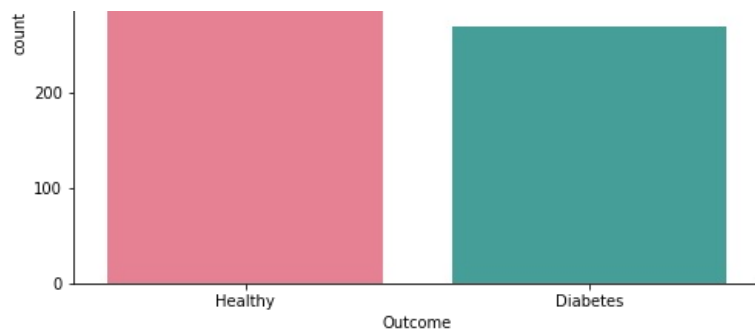
Out[14]: Text(0.5, 0, 'DiabetesPedigreeFunction')



```
In [15]: from matplotlib.pyplot import figure,show
figure(figsize=(8,6))
ax=sns.countplot(x=data['Outcome'],data=data,palette='husl')
ax.set_xticklabels(['Healthy','Diabetes'])
healthy,diabetes=data['Outcome'].value_counts().values
print('Samples of diabetes people:',diabetes)
print('Samples of healthy people: ',healthy)
```

Samples of diabetes people: 268
Samples of healthy people: 500






Distribution of other features with respect to Outcome


```
In [16]: plt.figure()
ax=sns.histplot(data['Pregnancies'][data.Outcome==1],color='turquoise',kde=True,stat='density',edgecolor='none')
sns.histplot(data['Pregnancies'][data.Outcome==0],color='coral',kde=True,stat='density',edgecolor='none')
plt.legend(['Diabetes','No Diabetes'])
```

Out[16]: <matplotlib.legend.Legend at 0x1de3167f640>



```
In [17]: plt.figure()
ax=sns.histplot(data['Glucose'][data.Outcome==1],color='darkturquoise',kde=True,stat='density',edgecolor='none')
sns.histplot(data['Glucose'][data.Outcome==0],color='lightcoral',kde=True,stat='density',edgecolor='none')
plt.legend(['Diabetes','No Diabetes'])
```


Out[17]: <matplotlib.legend.Legend at 0x1de31e430a0>



```
In [18]: plt.figure()
ax=sns.histplot(data['BloodPressure'][data.Outcome==1],color='darkturquoise',kde=True,stat='density',edgecolor='r')
sns.histplot(data['BloodPressure'][data.Outcome==0],color='lightcoral',kde=True,stat='density',edgecolor='none')
plt.legend(['Diabetes','No Diabetes'])
```


Out[18]: <matplotlib.legend.Legend at 0x1de32442100>






```
In [19]: plt.figure()
ax=sns.histplot(data['SkinThickness'][data.Outcome==1],color='darkturquoise',kde=True,stat='density',edgecolor='r')
sns.histplot(data['SkinThickness'][data.Outcome==0],color='lightcoral',kde=True,stat='density',edgecolor='none')
plt.legend(['Diabetes','No Diabetes'])
```

Out[19]: <matplotlib.legend.Legend at 0x1de324422b0>




```
In [20]: plt.figure()
ax=sns.histplot(data['Insulin'][data.Outcome==1],color='turquoise',kde=True,stat='density',edgecolor='none')
sns.histplot(data['Insulin'][data.Outcome==0],color='coral',kde=True,stat='density',edgecolor='none')
plt.legend(['Diabetes','No Diabetes'])
```


Out[20]: <matplotlib.legend.Legend at 0x1de335c4040>



```
In [21]: plt.figure()
ax=sns.histplot(data['DiabetesPedigreeFunction'][data.Outcome==1],color='darkturquoise',kde=True,stat='density',edgecolor='r')
sns.histplot(data['DiabetesPedigreeFunction'][data.Outcome==0],color='lightcoral',kde=True,stat='density',edgecolor='none')
plt.legend(['Diabetes','No Diabetes'])
```


Out[21]: <matplotlib.legend.Legend at 0x1de335a59d0>





```
In [22]: plt.figure()
ax=sns.histplot(data['Age'][data.Outcome==1],color='darkturquoise',kde=True,stat='density',edgecolor='none')
sns.histplot(data['Age'][data.Outcome==0],color='lightcoral',kde=True,stat='density',edgecolor='none')
sns.histplot(data['Age'],color='green',kde=True,stat='density',linewidth=0)
plt.legend(['Diabetes','No Diabetes','all'])
```

Out[22]: <matplotlib.legend.Legend at 0x1de324aad00>

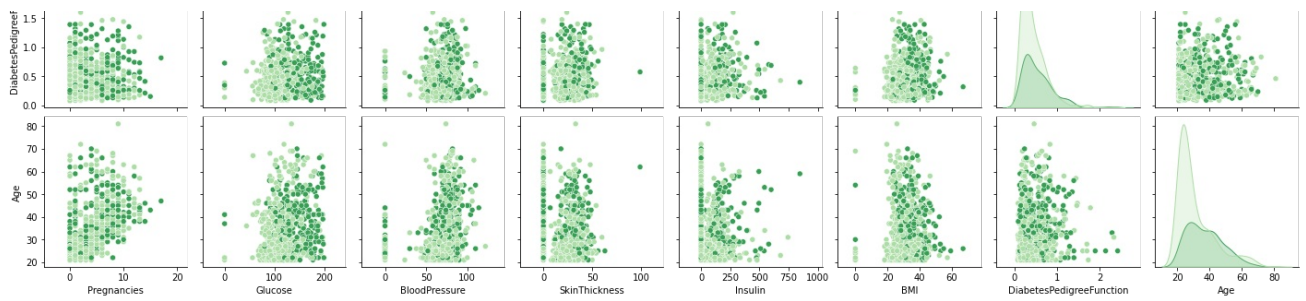


Pairplot

```
In [23]: sns.pairplot(data,hue='Outcome',palette='Greens')
```

Out[23]: <seaborn.axisgrid.PairGrid at 0x1de31e2c7c0>






Heatmap Plot Finding Correlation between the columns

```
In [24]: matrix=data.corr()
f,ax=plt.subplots(figsize=(18,18))
sns.heatmap(matrix,vmax=.7,square=True,cmap='YlGnBu',annot=True,linewidth=.5).set_title('Correlation Map')
```

```
Out[24]: Text(0.5, 1.0, 'Correlation Map')
```



Checking for values which are 0

```
In [25]: print('Total number of rows in the dataset:',len(data))
print('Total number of rows in the Glucose which are zero:',len(data.loc[data['Glucose']==0]))
print('Total number of rows in the Blood Pressure which are zero:',len(data.loc[data['BloodPressure']==0]))
print('Total number of rows in the Skin Thickness which are zero:',len(data.loc[data['SkinThickness']==0]))
print('Total number of rows in the Insulin which are zero:',len(data.loc[data['Insulin']==0]))
print('Total number of rows in the BMI which are zero:',len(data.loc[data['BMI']==0]))
```

```
print('Total number of rows in the DiabetesPedigreeFunction which are zero:',len(data.loc[data['DiabetesPedigreeFunction']==0]))
print('Total number of rows in the Age which are zero:',len(data.loc[data['Age']==0]))
```

```
Total number of rows in the dataset: 768
Total number of rows in the Glucose which are zero: 5
Total number of rows in the Blood Pressure which are zero: 35
Total number of rows in the Skin Thickness which are zero: 227
Total number of rows in the Insulin which are zero: 374
Total number of rows in the BMI which are zero: 11
Total number of rows in the DiabetesPedigreeFunction which are zero: 0
Total number of rows in the Age which are zero: 0
```

```
In [26]: # separating last column i.e outcomes from the dataset and assigning to different variables
data['Outcome'].value_counts()
```

```
Out[26]: Outcome
0      500
1      268
Name: count, dtype: int64
```

```
In [27]: data.groupby('Outcome').mean()
```

```
Out[27]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age
Outcome								
0	3.298000	109.980000	68.184000	19.664000	68.792000	30.304200	0.429734	31.190000
1	4.865672	141.257463	70.824627	22.164179	100.335821	35.142537	0.550500	37.067164

Separating data and labels

```
In [28]: x=data.drop(columns='Outcome',axis=1)
y=data['Outcome']
```

```
In [29]: print(x)
```

```

Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin  BMI  \
0             6      148             72             35         0  33.6
1             1       85             66             29         0  26.6
2             8      183             64              0         0  23.3
3             1       89             66             23        94  28.1
4             0      137             40             35       168  43.1
..          ...      ...             ...             ...      ...  ...
763          10      101             76             48       180  32.9
764           2      122             70             27         0  36.8
765           5      121             72             23       112  26.2
766           1      126             60              0         0  30.1
767           1       93             70             31         0  30.4

DiabetesPedigreeFunction  Age
0                      0.627  50
1                      0.351  31
2                      0.672  32
3                      0.167  21
4                      2.288  33
..                      ...   ...
763                    0.171  63
764                    0.340  27
765                    0.245  30
766                    0.349  47
767                    0.315  23
```

```
[768 rows x 8 columns]
```

```
In [30]: print(y)
```

```
0      1
1      0
2      1
3      0
```

```
4      1
      ..
763    0
764    0
765    0
766    1
767    0
Name: Outcome, Length: 768, dtype: int64
```

Data Standardisation

```
In [31]: scaler=StandardScaler()
        scaler.fit(x)
```

```
Out[31]: StandardScaler()
```

```
In [32]: standardised_data=scaler.transform(x)
        print(standardised_data)

[[ 0.63994726  0.84832379  0.14964075 ...  0.20401277  0.46849198
    1.4259954 ]
 [-0.84488505 -1.12339636 -0.16054575 ... -0.68442195 -0.36506078
   -0.19067191]
 [ 1.23388019  1.94372388 -0.26394125 ... -1.10325546  0.60439732
   -0.10558415]
 ...
 [ 0.3429808  0.00330087  0.14964075 ... -0.73518964 -0.68519336
   -0.27575966]
 [-0.84488505  0.1597866  -0.47073225 ... -0.24020459 -0.37110101
    1.17073215]
 [-0.84488505 -0.8730192   0.04624525 ... -0.20212881 -0.47378505
   -0.87137393]]
```

```
In [33]: x=standardised_data
        y=data['Outcome']
        print(x)
        print(y)

[[ 0.63994726  0.84832379  0.14964075 ...  0.20401277  0.46849198
    1.4259954 ]
 [-0.84488505 -1.12339636 -0.16054575 ... -0.68442195 -0.36506078
   -0.19067191]
 [ 1.23388019  1.94372388 -0.26394125 ... -1.10325546  0.60439732
   -0.10558415]
 ...
 [ 0.3429808  0.00330087  0.14964075 ... -0.73518964 -0.68519336
   -0.27575966]
 [-0.84488505  0.1597866  -0.47073225 ... -0.24020459 -0.37110101
    1.17073215]
 [-0.84488505 -0.8730192   0.04624525 ... -0.20212881 -0.47378505
   -0.87137393]]
0      1
1      0
2      1
3      0
4      1
      ..
763    0
764    0
765    0
766    1
767    0
Name: Outcome, Length: 768, dtype: int64
```

Splitting data into train data and test data

```
In [34]: x_train,x_test,y_train,y_test = train_test_split(x,y, test_size= 0.3, stratify=y,random_state=2)
        print(x.shape,x_train.shape,x_test.shape)
```

```
(768, 8) (537, 8) (231, 8)
```

Training the model

```
In [35]: classifier = svm.SVC(kernel='linear')
```

Training the support vector Machine Classifier

```
In [36]: classifier.fit(x_train,y_train)
```

```
Out[36]: SVC(kernel='linear')
```

Accuracy score on the training data

```
In [37]: x_train_prediction = classifier.predict(x_train)
training_data_accuracy=accuracy_score(x_train_prediction,y_train)
```

```
In [38]: print('Accuracy score of the training data:',training_data_accuracy)
```

Accuracy score of the training data: 0.7821229050279329

Making a predictive system

```
In [39]: input_data=(1,103,30,38,83,43.3,0.183,33) #giving input from dataset and checking whether the prediction is right
input_data_as_numpy_array=np.asarray(input_data) #changing the input data to numpy array
input_data_resaped=input_data_as_numpy_array.reshape(1,-1) # reshaping the array as we are predicting only for c
std_data=scaler.transform(input_data_resaped) # standardising the input data
print(std_data)

prediction = classifier.predict(std_data)
print(prediction)

if (prediction[0]==0):
    print('The person is NOT DIABETIC')
else:
    print('The person is DIABETIC')
```

[[-0.84488505 -0.56004775 -2.02166474 1.09545411 0.02778979 1.43512945
 -0.87244072 -0.0204964]]

[0]
The person is NOT DIABETIC

```
In [40]: # represents a linear stack of layers
from keras.models import Sequential
# fully connected layer
from keras.layers import Dense

# Model Evaluation
from sklearn import metrics
from sklearn.metrics import classification_report, confusion_matrix, ConfusionMatrixDisplay, recall_score, accuracy_score
import pickle
```

```
In [41]: # Normality Test
from scipy import stats
## retrieving p value from normality test function
PregnanciesPVAL=stats.normaltest(data.Pregnancies).pvalue
GlucosePVAL=stats.normaltest(data.Glucose).pvalue
BloodPressurePVAL=stats.normaltest(data.BloodPressure).pvalue
SkinThicknessPVAL=stats.normaltest(data.SkinThickness).pvalue
InsulinPVAL=stats.normaltest(data.Insulin).pvalue
BMIPVAL=stats.normaltest(data.BMI).pvalue
diaPeFuPVAL=stats.normaltest(data.DiabetesPedigreeFunction).pvalue
AgePVAL=stats.normaltest(data.Age).pvalue
## Printing the values
print("Pregnancies P Value is " + str(PregnanciesPVAL))
print("Glucose P Value is " + str(GlucosePVAL))
print("BloodPressure P Value is " + str(BloodPressurePVAL))
```

```
print("Skin Thickness P Value is " + str(SkinThicknessPVAL))
print("Insulin P Value is " + str(InsulinPVAL))
print("BMI P Value is " + str(BMIPVAL))
print("dfbetes Pedigree Function P Value is " + str(diaPeFuPVAL))
print("Age P Value is " + str(AgePVAL))
```

```
Pregnancies P Value is 3.9142916387836715e-18
Glucose P Value is 0.0020446506991363502
BloodPressure P Value is 3.780127075543807e-67
Skin Thickness P Value is 0.00017124157678801767
Insulin P Value is 6.895342740914799e-85
BMI P Value is 1.9696869549260572e-19
dfbetes Pedigree Function P Value is 1.2987697487612696e-70
Age P Value is 9.308980038236578e-27
```

Conclusion

None of the variables are normal. ($P > 0.05$)

The variables "Pregnancies," "Glucose," "Insulin," "BMI," "Diabetes Pedigree Function," and "Age" show strong evidence of association with the outcome variable based on their small p-values. The variables "Blood Pressure" and "Skin Thickness" also show some evidence of association, although the evidence may not be as strong.

In [42]:

```
import plotly.graph_objs as go
textd = ['non-diabetic' if cl==0 else 'diabetic' for cl in data['Outcome']]

fig = go.Figure(data=go.Splom(
    dimensions=[dict(label='Pregnancies', values=data['Pregnancies']),
                  dict(label='Glucose', values=data['Glucose']),
                  dict(label='BloodPressure', values=data['BloodPressure']),
                  dict(label='SkinThickness', values=data['SkinThickness']),
                  dict(label='Insulin', values=data['Insulin']),
                  dict(label='BMI', values=data['BMI']),
                  dict(label='DiabPedigreeFun', values=data['DiabetesPedigreeFunction']),
                  dict(label='Age', values=data['Age'])],
    showupperhalf=False, # remove plots on diagonal
    marker=dict(color=data['Outcome'],
                size=5,
                colorscale='Viridis',
                line=dict(width=0.5,
                        color='rgb(230,230,230)')),
    text=textd))
#diagonal=dict(visible=False))

title = "Scatterplot Matrix for Diabetes Dataset"
fig.update_layout(title=title,
                  template = 'plotly_white',
                  dragmode='select',
                  width=1000,
                  height=1000,
                  hovermode='closest')

fig.show()
```

info

Note: In the above boxplots yellow points refers Diabetic and dark blue points refers Non-diabetic

Conclusion

Only BMI & SkinThickness and Pregnancies & Age seem to have positive linear relationships. Another likely suspect is Glucose and Insulin. There are no non-linear relationships.

Thank You

In []: